

```
#from google.colab import drive
#drive.mount('/content/drive')

import os
#os.chdir('/content/drive/MyDrive')

#os.environ['KAGGLE_CONFIG_DIR'] = '/content/drive/MyDrive'

#! kaggle datasets download -d yacharki/yelp-reviews-for-sentianalysis-binary-np-csv

#!kaggle datasets download -d yacharki/yelp-reviews-for-sentianalysis-binary-np-csv -p /content

#!mv /content/yelp-reviews-for-sentianalysis-binary-np-csv.zip /content/drive/MyDrive/

#!unzip /content/drive/MyDrive/yelp-reviews-for-sentianalysis-binary-np-csv.zip -d /content/drive/MyDrive/

CUDA_LAUNCH_BLOCKING=1
import numpy as np
import pandas as pd
import re
import nltk
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer,TfidfVectorizer
from sklearn import preprocessing,tree
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score,recall_score,confusion_matrix,f1_score
from sklearn.model_selection import GridSearchCV
from scipy import sparse
from sklearn.svm import LinearSVC
from sklearn.naive_bayes import MultinomialNB
#from xgboost import XGBClassifier
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, accuracy_score
from sklearn.pipeline import Pipeline
import matplotlib.pyplot as plt

yelp_train_data = pd.read_csv('/content/drive/MyDrive/yelp_review_sa_binary_csv/train.csv')
yelp_test_data = pd.read_csv('/content/drive/MyDrive/yelp_review_sa_binary_csv/test.csv')
```

```
print(yelp_train_data.shape)
print(yelp_test_data.shape)
```

```
(56000, 2)
(38000, 2)
```

## Text Cleaning

```
print(yelp_train_data.isnull().sum())
print(yelp_test_data.isnull().sum())
```

```
class_index    0
review_text    0
dtype: int64
class_index    0
review_text    0
dtype: int64
```

```
yelp_train_data = yelp_train_data[yelp_train_data['review_text'].notna()]
yelp_train_data = yelp_train_data[yelp_train_data['review_text'].str.strip() != ""]
```

```
# Get the class distribution in the original data
class_counts = yelp_train_data['class_index'].value_counts(normalize=True)
```

```
# Calculate the number of samples per class for a total of 5,000 rows
samples_per_class = (class_counts * 5000).round().astype(int)
```

```
# Perform stratified sampling
sampled_data = yelp_train_data.groupby('class_index').apply(lambda x: x.sample(n=samples_per_class[x.name], random_state=42)).reset_index()
```

```
# Check the class distribution in the sampled data
print(sampled_data['class_index'].value_counts(normalize=True))
```

```
↗ class_index
1    0.5
2    0.5
Name: proportion, dtype: float64
<ipython-input-23-6999dce83965>:11: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated. Use .apply() instead.
sampled_data = yelp_train_data.groupby('class_index').apply(lambda x: x.sample(n=samples_per_class[x.name], random_state=42)).reset_index()
```

```
sampled_data.to_csv("yelp_train_data.csv", index=False)
```

```
yelp_test_data = yelp_test_data[yelp_test_data['review_text'].notna()]
yelp_test_data = yelp_test_data[yelp_test_data['review_text'].str.strip() != ""]
```

```
# Get the class distribution in the original data
class_counts = yelp_test_data['class_index'].value_counts(normalize=True)
```

```
# Calculate the number of samples per class for a total of 1000 rows
samples_per_class = (class_counts * 1000).round().astype(int)
```

```
# Perform stratified sampling
sampled_data = yelp_test_data.groupby('class_index').apply(lambda x: x.sample(n=samples_per_class[x.name], random_state=42)).reset_index()
```

```
# Check the class distribution in the sampled data
print(sampled_data['class_index'].value_counts(normalize=True))
```

```
↗ class_index
1    0.5
2    0.5
Name: proportion, dtype: float64
<ipython-input-25-22511f65c20d>:11: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated. Use .apply() instead.
sampled_data = yelp_test_data.groupby('class_index').apply(lambda x: x.sample(n=samples_per_class[x.name], random_state=42)).reset_index()
```

```
sampled_data.to_csv("yelp_test_data.csv", index=False)
```

```
def cleaning_text(review):
```

```
    #removing the url's
    review = re.sub('http\S+\s*', ' ', review)
    #removing the punctuations
    review = re.sub('[%s]' % re.escape("""!"#$%&'()*+,-./:;<=>@[\\]^_`{|}~?"""), ' ', review)
    #removing non-ascii characters
    review = re.sub(r'^\x00-\x7f', r' ', review)
    #removing mentions (i.e, @)
    review = re.sub('@\S+', ' ', review)
    #removing hashtags
    review = re.sub('#\S+', ' ', review)
    #remove numbers
    review = re.sub("\d+", ' ', review)
    #removing extra whitespaces, wherever applicable
    review = re.sub('\s+', ' ',review)
    #converting the text into lowercase
    review = review.lower()
```

```
    return review
```

```
yelp_train_data = pd.read_csv('yelp_train_data.csv')
yelp_test_data = pd.read_csv('yelp_test_data.csv')
```

```
yelp_train_data['clean_review_text'] = yelp_train_data['review_text'].apply(cleaning_text)
yelp_test_data['clean_review_text'] = yelp_test_data['review_text'].apply(cleaning_text)
```

```
yelp_train_data.head(5)
```

```
↗
```

|   | class_index | review_text                                       | clean_review_text                                 |
|---|-------------|---|---|
| 0 | 1           | I do enjoy Taco Bell from time to time but I h... | i do enjoy taco bell from time to time but i h... |
| 1 | 1           | My sister and I are pretty great guest at rest... | my sister and i are pretty great guest at rest... |
| 2 | 1           | Le vrai four \u00e0 bois \u00e0 l'entr\u00e9e ... | le vrai four u e bois u e l entr u e e u e tai... |
| 3 | 1           | Felt discriminated because I came there in a s... | felt discriminated because i came there in a s... |
| 4 | 1           | Absolutely AWFUL service. \nWe were originallv... | absolutely awful service nwe were originallv t... |

```
yelp_test_data.head(5)
```

|   | class_index | review_text  | clean_review_text                                 |
|---|-------------|--|---|
| 0 | 1           | Worst Starbucks I've ever been to! The staff is... | worst starbucks i ve ever been to the staff is... |
| 1 | 1           | F*ck this place. Maybe the location in Tempe is... | f ck this place maybe the location in tempe is... |
| 2 | 1           | Stayed at the Hakone Suite, room is fantastic ...  | stayd at the hakone suite room is fantastic a...  |
| 3 | 1           | This place sucks. It's small, and their adver...   | this place sucks it s small and their advertis... |
| 4 | 1           | Ate there for the first and last time. Let me...   | ate there for the first and last time let me s... |

```
yelp_train_data.class_index.value_counts()
```

|             | count |
|-------------|-------|
| class_index |       |
| 1           | 2500  |
| 2           | 2500  |

```
yelp_test_data.class_index.value_counts()
```

|             | count |
|-------------|-------|
| class_index |       |
| 1           | 500   |
| 2           | 500   |

Tokenization

```
import nltk
nltk.download('punkt_tab')
```

```
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt_tab.zip.
True
```

```
yelp_train_data['tokens'] = yelp_train_data['clean_review_text'].apply(word_tokenize)
yelp_test_data['tokens'] = yelp_test_data['clean_review_text'].apply(word_tokenize)
```

```
yelp_train_data.head()
```

|   | class_index | review_text                                       | clean_review_text                                 | tokens  |
|---|-------------|---|---|---|
| 0 | 1           | I do enjoy Taco Bell from time to time but I h... | i do enjoy taco bell from time to time but i h... | [i, do, enjoy, taco, bell, from, time, to, tim... |
| 1 | 1           | My sister and I are pretty great guest at rest... | my sister and i are pretty great guest at rest... | [my, sister, and, i, are, pretty, great, guest... |
| 2 | 1           | Le vrai four \u00e0 bois \u00e0 l'entr\u00e9e ... | le vrai four u e bois u e l entr u e e u e tai... | [le, vrai, four, u, e, bois, u, e, l, entr, u,... |

Felt discriminated because I came there in a ... felt discriminated because i came there in a ... felt discriminated, because i came

Remove Stopwords

```
nltk.download('stopwords')
stop_words = set(stopwords.words('english'))
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
```

```
yelp_train_data['tokens'] = yelp_train_data['tokens'].apply(lambda x: [word for word in x if word not in stop_words])
yelp_test_data['tokens'] = yelp_test_data['tokens'].apply(lambda x: [word for word in x if word not in stop_words])
```

```
yelp_train_data.head()
```

|   | class_index | review_text                                       | clean_review_text                                 | tokens  |
|---|-------------|---|---|---|
| 0 | 1           | I do enjoy Taco Bell from time to time but I h... | i do enjoy taco bell from time to time but i h... | [enjoy, taco, bell, time, time, hate, say, las... |
| 1 | 1           | My sister and I are pretty great guest at rest... | my sister and i are pretty great guest at rest... | [sister, pretty, great, guest, restaurants, ti... |
| 2 | 1           | Le vrai four \u00e0 bois \u00e0 l'entr\u00e9e ... | le vrai four u e bois u e l entr u e e u e tai... | [le, vrai, four, u, e, bois, u, e, l, entr, u,... |

```
yelp_test_data.head()
```

|   | class_index | review_text                                       | clean_review_text                                 | tokens  |
|---|-------------|---|---|---|
| 0 | 1           | Worst Starbucks I've ever been to! The staff i... | worst starbucks i ve ever been to the staff is... | [worst, starbucks, ever, staff, extremely, rud... |
| 1 | 1           | F*ck this place. Maybe the location in Tempe l... | f ck this place maybe the location in tempe is... | [f, ck, place, maybe, location, tempe, better,... |
| 2 | 1           | Stayed at the Hakone Suite, room is fantastic ... | stayed at the hakone suite room is fantastic a... | [stayed, hakone, suite, room, fantastic, satis... |

## Lemmatization

```
nlTK.download('wordnet')
nlTK.download('omw-1.4')
```

```
[nlTK_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Downloading package omw-1.4 to /root/nltk_data...
True
```

```
lemmatizer = WordNetLemmatizer()
```

```
yelp_train_data['tokens'] = yelp_train_data['tokens'].apply(lambda x: [lemmatizer.lemmatize(word) for word in x])
yelp_test_data['tokens'] = yelp_test_data['tokens'].apply(lambda x: [lemmatizer.lemmatize(word) for word in x])
```

```
yelp_train_data.head()
```

|   | class_index | review_text                                       | clean_review_text                                 | tokens  |
|---|-------------|---|---|---|
| 0 | 1           | I do enjoy Taco Bell from time to time but I h... | i do enjoy taco bell from time to time but i h... | [enjoy, taco, bell, time, time, hate, say, las... |
| 1 | 1           | My sister and I are pretty great guest at rest... | my sister and i are pretty great guest at rest... | [sister, pretty, great, guest, restaurant, tip... |
| 2 | 1           | Le vrai four \u00e0 bois \u00e0 l'entr\u00e9e ... | le vrai four u e bois u e l entr u e e u e tai... | [le, vrai, four, u, e, bois, u, e, l, entr, u,... |

```
yelp_train_data['processed_review'] = yelp_train_data['tokens'].apply(lambda x: ' '.join(x))
yelp_test_data['processed_review'] = yelp_test_data['tokens'].apply(lambda x: ' '.join(x))
```

```
yelp_train_data.head()
```

|   | class_index | review_text                                       | clean_review_text                                 | tokens  | processed_review                                  |
|---|-------------|---|---|---|---|
| 0 | 1           | I do enjoy Taco Bell from time to time but I h... | i do enjoy taco bell from time to time but i h... | [enjoy, taco, bell, time, time, hate, say, las... | enjoy taco bell time time hate say last time f... |
| 1 | 1           | My sister and I are pretty great guest at rest... | my sister and i are pretty great guest at rest... | [sister, pretty, great, guest, restaurant, tip... | sister pretty great guest restaurant tip well ... |
| 2 | 1           | Le vrai four \u00e0 bois \u00e0 l'entr\u00e9e ... | le vrai four u e bois u e l entr u e e u e tai... | [le, vrai, four, u, e, bois, u, e, l, entr, u,... | le vrai four u e bois u e l entr u e e u e tai... |

```
X_train = yelp_train_data['processed_review'].tolist()
y_train = yelp_train_data['class_index'].tolist()
```

```
X_test = yelp_test_data['processed_review'].tolist()
y_test = yelp_test_data['class_index'].tolist()
```

```
label_mapping = {label: idx for idx, label in enumerate(sorted(set(y_train)))}
y_train = [label_mapping[label] for label in y_train]
```

```
label_mapping = {label: idx for idx, label in enumerate(sorted(set(y_test)))}
y_test = [label_mapping[label] for label in y_test]
```

```
from transformers import BertTokenizer, BertForSequenceClassification, AdamW
from torch.utils.data import DataLoader, Dataset
import torch
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

```
tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
```

```
def tokenize_data(texts, labels, max_length=128):
    tokenized = tokenizer(
        texts,
        max_length=max_length,
        padding="max_length",
        truncation=True,
        return_tensors="pt"
    )
    return tokenized["input_ids"], tokenized["attention_mask"], torch.tensor(labels)
```

```
input_ids, attention_masks, labels = tokenize_data(X_train, y_train)
```

⚡ /usr/local/lib/python3.10/dist-packages/huggingface\_hub/utils/\_auth.py:94: UserWarning:  
The secret `HF\_TOKEN` does not exist in your Colab secrets.  
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as :  
You will be able to reuse this secret in all of your notebooks.  
Please note that authentication is recommended but still optional to access public models or datasets.

```
warnings.warn(
tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 1.73kB/s]
vocab.txt: 100% 232k/232k [00:00<00:00, 503kB/s]
tokenizer.json: 100% 466k/466k [00:00<00:00, 736kB/s]
config.json: 100% 570/570 [00:00<00:00, 41.2kB/s]
```

```
class TextDataset(Dataset):
    def __init__(self, input_ids, attention_masks, labels):
        self.input_ids = input_ids
        self.attention_masks = attention_masks
        self.labels = labels

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, idx):
        return {
            'input_ids': self.input_ids[idx],
            'attention_mask': self.attention_masks[idx],
            'labels': self.labels[idx]
        }
```

```
X_train_ids, X_val_ids, train_masks, val_masks, y_train, y_val = train_test_split(
    input_ids, attention_masks, labels, test_size=0.2, random_state=42
)
```

```
train_dataset = TextDataset(X_train_ids, train_masks, y_train)
val_dataset = TextDataset(X_val_ids, val_masks, y_val)
```

```
train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=8)
```

```
torch.cuda.empty_cache()
%env CUDA_LAUNCH_BLOCKING=1
```

⚡ env: CUDA\_LAUNCH\_BLOCKING=1

```
torch.cuda.is_available()
```

⚡ True

```
model = BertForSequenceClassification.from_pretrained("bert-base-uncased", num_labels=2)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
```

```
# Optimizer
optimizer = AdamW(model.parameters(), lr=2e-5)
```

 Show hidden output

```
# Early stopping parameters
early_stopping_patience = 2
best_val_loss = float("inf")
patience_counter = 0
```

```
# Lists to store metrics for plotting
train_losses = []
val_losses = []
val_accuracies = []
```

```
epochs = 20
best_val_accuracy = 0 # Initialize for early stopping
```

```
for epoch in range(epochs):
```

```
    # Training Phase
```

```
    model.train()
    total_loss = 0
    train_predictions = []
    train_labels = []
```

```
    for batch in train_loader:
        optimizer.zero_grad()
```

```
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['labels'].to(device)
```

```
        outputs = model(input_ids=input_ids, attention_mask=attention_mask, labels=labels)
        loss = outputs.loss
        total_loss += loss.item()
```

```
        # Track training predictions and labels
        preds = torch.argmax(outputs.logits, dim=1)
        train_predictions.extend(preds.cpu().numpy())
        train_labels.extend(labels.cpu().numpy())
```

```
        loss.backward()
        optimizer.step()
```

```
    avg_train_loss = total_loss / len(train_loader)
    train_losses.append(avg_train_loss)
```

```
    train_accuracy = accuracy_score(train_labels, train_predictions)
    print(f"Epoch {epoch + 1}, Training Loss: {avg_train_loss:.4f}, Training Accuracy: {train_accuracy:.4f}")
```

```
    # Validation Phase
```

```
    model.eval()
    val_loss = 0
    val_predictions = []
    val_labels = []
```

```
    with torch.no_grad():
```

```
        for batch in val_loader:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['labels'].to(device)
```

```
            outputs = model(input_ids=input_ids, attention_mask=attention_mask, labels=labels)
            val_loss += outputs.loss.item()
```

```
            preds = torch.argmax(outputs.logits, dim=1)
            val_predictions.extend(preds.cpu().numpy())
            val_labels.extend(labels.cpu().numpy())
```

```
    avg_val_loss = val_loss / len(val_loader)
    val_losses.append(avg_val_loss)
```

```
    val_accuracy = accuracy_score(val_labels, val_predictions)
    val_accuracies.append(val_accuracy)
```

```

print(f"Epoch {epoch + 1}, Validation Loss: {avg_val_loss:.4f}, Validation Accuracy: {val_accuracy:.4f}")

# Early Stopping Check
if avg_val_loss < best_val_loss or val_accuracy > best_val_accuracy:
    best_val_loss = avg_val_loss
    best_val_accuracy = val_accuracy
    patience_counter = 0 # Reset patience counter
else:
    patience_counter += 1
    print(f"No improvement for {patience_counter} epoch(s).")

if patience_counter >= early_stopping_patience:
    print("Early stopping triggered.")
    break

```

```

Epoch 1, Training Loss: 0.3651, Training Accuracy: 0.8365
Epoch 1, Validation Loss: 0.3097, Validation Accuracy: 0.8670
Epoch 2, Training Loss: 0.1908, Training Accuracy: 0.9267
Epoch 2, Validation Loss: 0.3295, Validation Accuracy: 0.8710
Epoch 3, Training Loss: 0.1010, Training Accuracy: 0.9647
Epoch 3, Validation Loss: 0.3925, Validation Accuracy: 0.8910
Epoch 4, Training Loss: 0.0406, Training Accuracy: 0.9875
Epoch 4, Validation Loss: 0.3416, Validation Accuracy: 0.8920
Epoch 5, Training Loss: 0.0452, Training Accuracy: 0.9855
Epoch 5, Validation Loss: 0.4495, Validation Accuracy: 0.8870
No improvement for 1 epoch(s).
Epoch 6, Training Loss: 0.0294, Training Accuracy: 0.9892
Epoch 6, Validation Loss: 0.4842, Validation Accuracy: 0.8810
No improvement for 2 epoch(s).
Early stopping triggered.

```

```

plt.figure(figsize=(10, 6))
plt.plot(range(1, len(train_losses) + 1), train_losses, label="Training Loss", marker='o')
plt.plot(range(1, len(val_losses) + 1), val_losses, label="Validation Loss", marker='o')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training and Validation Loss")
plt.legend()
plt.grid(True)
plt.show()

```



```

model.save_pretrained("bert-finetuned-model")
tokenizer.save_pretrained("bert-finetuned-model")

('bert-finetuned-model/tokenizer_config.json',
 'bert-finetuned-model/special_tokens_map.json',
 'bert-finetuned-model/vocab.txt',
 'bert-finetuned-model/added_tokens.json')

```

```
input_ids, attention_mask, labels = tokenize_data(X_test, y_test)
```

```

model.eval()
with torch.no_grad():
    input_ids = input_ids.to(device)
    attention_mask = attention_mask.to(device)
    outputs = model(input_ids=input_ids, attention_mask=attention_mask)
    predictions = torch.argmax(outputs.logits, dim=1).cpu().numpy()

print("Classification Report:")
print(classification_report(y_test, predictions, target_names=["Positive - class 1", "Negative - class 0"]))

```

```

↗ Classification Report:

```

|                    | precision | recall | f1-score | support |
|--------------------|-----------|--------|----------|---------|
| Positive - class 1 | 0.87      | 0.93   | 0.90     | 500     |
| Negative - class 0 | 0.92      | 0.86   | 0.89     | 500     |
| accuracy           |           |        | 0.90     | 1000    |
| macro avg          | 0.90      | 0.90   | 0.89     | 1000    |
| weighted avg       | 0.90      | 0.90   | 0.89     | 1000    |

```

from sklearn.metrics import roc_curve, auc

fpr, tpr, thresholds = roc_curve(y_test, predictions)
roc_auc = auc(fpr, tpr)

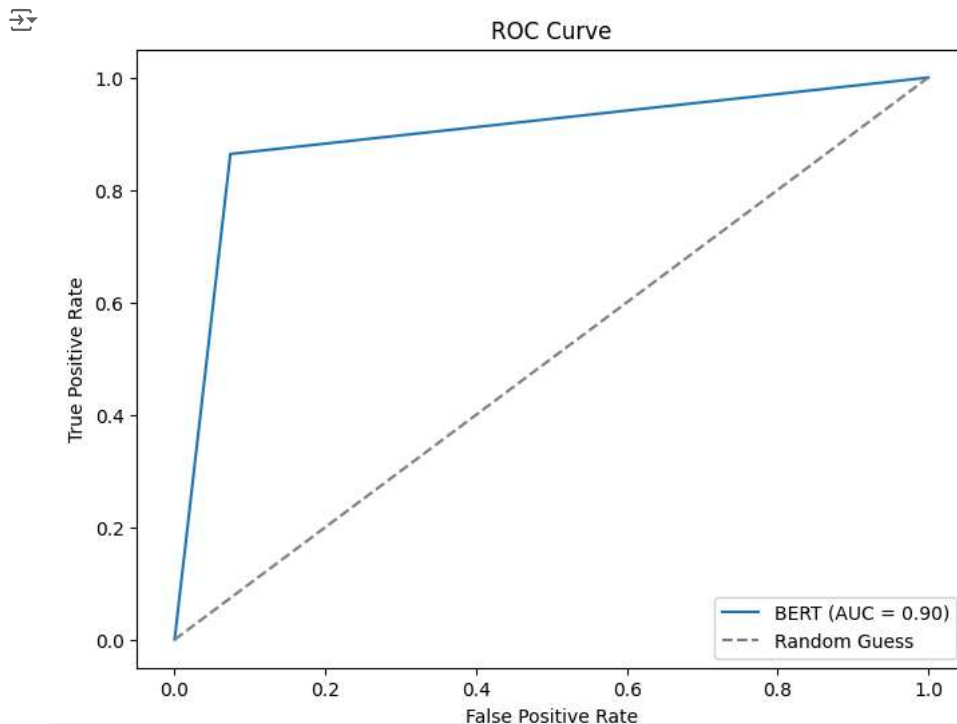
```

```
import matplotlib.pyplot as plt
```

```

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'BERT (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()

```



```

import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

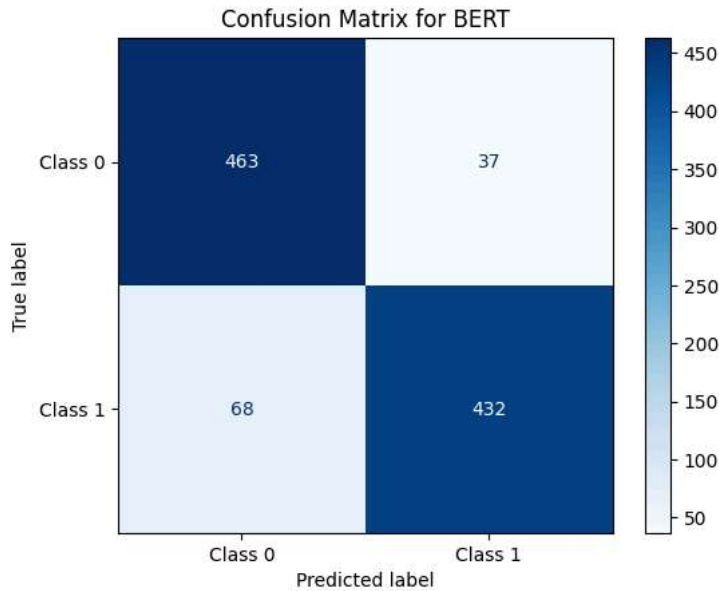
cm_bert = confusion_matrix(y_test, predictions)
disp_bert = ConfusionMatrixDisplay(confusion_matrix=cm_bert, display_labels=["Class 0", "Class 1"])

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
disp_bert.plot(cmap="Blues", values_format="d")
plt.title("Confusion Matrix for BERT")
plt.show()

```



&lt;Figure size 800x600 with 0 Axes&gt;



```

X_train = yelp_train_data['processed_review'].tolist()
y_train = yelp_train_data['class_index'].tolist()

X_test = yelp_test_data['processed_review'].tolist()
y_test = yelp_test_data['class_index'].tolist()

label_mapping = {label: idx for idx, label in enumerate(sorted(set(y_train)))}
y_train = [label_mapping[label] for label in y_train]

label_mapping = {label: idx for idx, label in enumerate(sorted(set(y_test)))}
y_test = [label_mapping[label] for label in y_test]

def tf_idf_vect(train_data, test_data):

    #Initialising the tf-idf vectorizer
    #df: When building the vocabulary it ignores the terms that have a document frequency strictly lower than the given threshold(df=5)
    vectorizer = TfidfVectorizer(min_df=8)
    #fit_transform learns the vocabulary dictionary and return document-term matrix.
    train_cv = vectorizer.fit_transform(train_data)
    #Transform documents to document-term matrix
    test_cv = vectorizer.transform(test_data)

    return train_cv, test_cv

X_train, X_test = tf_idf_vect(X_train, X_test)

# Define the pipeline with Multinomial Naive Bayes
pipeline_nb = Pipeline([
    ('nb', MultinomialNB())
])

# Define the parameter grid for GridSearchCV
param_grid_nb = {
    'nb__alpha': [0.01, 0.1, 1, 10] # Smoothing parameter for Naive Bayes
}

# Setup GridSearchCV
grid_nb = GridSearchCV(pipeline_nb, param_grid_nb, cv=5, scoring='accuracy', verbose=1)

# Train the model with GridSearchCV on the training data
grid_nb.fit(X_train, y_train) # assuming y_train contains the labels for training

# Best model, predict, and evaluate
best_nb = grid_nb.best_estimator_
y_pred_nb = best_nb.predict(X_test)
y_prob_nb = best_nb.predict_proba(X_test)[:, 1] # Probability scores for ROC curve

# Display best parameters and cross-validation score
print("Best parameters:", grid_nb.best_params_)
print("Best cross-validation score: {:.2f}".format(grid_nb.best_score_))

# Print classification report
print("\nClassification Report:")

```

```
print(classification_report(y_test, y_pred_nb)) # assuming y_test contains the labels for testing
```

```
# Print ROC-AUC score
```

```
print("ROC-AUC Score:", roc_auc_score(y_test, y_prob_nb))
```

```
# Plot ROC curve
```

```
fpr_nb, tpr_nb, _ = roc_curve(y_test, y_prob_nb, pos_label=1)
```

```
plt.figure(figsize=(8, 6))
```

```
plt.plot(fpr_nb, tpr_nb, label='Naive Bayes (AUC = %0.2f)' % roc_auc_score(y_test, y_prob_nb))
```

```
plt.plot([0, 1], [0, 1], 'r--')
```

```
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
```

```
plt.title('ROC Curve for Naive Bayes')
```

```
plt.legend(loc="lower right")
```

```
plt.show()
```

↗ Fitting 5 folds for each of 4 candidates, totalling 20 fits


Best parameters: {'nb\_\_alpha': 10}

Best cross-validation score: 0.86

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.83      | 0.85   | 0.84     | 500     |
| 1            | 0.85      | 0.83   | 0.84     | 500     |
| accuracy     |           |        | 0.84     | 1000    |
| macro avg    | 0.84      | 0.84   | 0.84     | 1000    |
| weighted avg | 0.84      | 0.84   | 0.84     | 1000    |

ROC-AUC Score: 0.924264



```
cm_nb = confusion_matrix(y_test, y_pred_nb)
```

```
disp_nb = ConfusionMatrixDisplay(confusion_matrix=cm_nb, display_labels=["Class 0", "Class 1"])
```

```
# Plot the confusion matrix
```


```
plt.figure(figsize=(8, 6))
```

```
disp_nb.plot(cmap="Blues", values_format="d")
```

```
plt.title("Confusion Matrix for Naive Bayes")
```

```
plt.show()
```

<Figure size 800x600 with 0 Axes>



```

from sklearn.tree import DecisionTreeClassifier
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import classification_report, roc_auc_score, roc_curve
import matplotlib.pyplot as plt

# Define the pipeline with Decision Tree
pipeline_dt = Pipeline([
    ('dt', DecisionTreeClassifier()) # Basic Decision Tree Classifier
])

# Define the parameter grid for GridSearchCV
param_grid_dt = {
    'dt__max_depth': [3, 5, 10],
    'dt__min_samples_split': [2, 5], # Minimum samples required to split an internal node
    'dt__min_samples_leaf': [1, 2, 4], # Minimum samples required to be at a leaf node
}

# Setup GridSearchCV
grid_dt = GridSearchCV(pipeline_dt, param_grid_dt, cv=5, scoring='accuracy', verbose=1)

# Train the model with GridSearchCV on the training data
grid_dt.fit(X_train, y_train) # Assuming X_train and y_train are defined

# Get the best model
best_dt = grid_dt.best_estimator_

# Predict on the test set
y_pred_dt = best_dt.predict(X_test) # Assuming X_test is defined
y_prob_dt = best_dt.predict_proba(X_test)[:, 1] # Probability scores for the positive class (1)

# Display the best parameters and cross-validation score
print("Best parameters:", grid_dt.best_params_)
print("Best cross-validation score: {:.2f}".format(grid_dt.best_score_))

# Fitting 5 folds for each of 18 candidates, totalling 90 fits
# Best parameters: {'dt__max_depth': 10, 'dt__min_samples_leaf': 2, 'dt__min_samples_split': 2}
# Best cross-validation score: 0.74

# Print classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred_dt))

# Compute ROC-AUC score
roc_auc = roc_auc_score(y_test, y_prob_dt)
print("ROC-AUC Score:", roc_auc)

# Compute the ROC curve
fpr_dt, tpr_dt, _ = roc_curve(y_test, y_prob_dt, pos_label=1)

# Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_dt, tpr_dt, label='Decision Tree (AUC = %.2f)' % roc_auc)

```


```
plt.plot([0, 1], [0, 1], 'r--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Decision Tree')
plt.legend(loc="lower right")
plt.show()
```



Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.68      | 0.83   | 0.75     | 500     |
| 1            | 0.78      | 0.61   | 0.69     | 500     |
| accuracy     |           |        | 0.72     | 1000    |
| macro avg    | 0.73      | 0.72   | 0.72     | 1000    |
| weighted avg | 0.73      | 0.72   | 0.72     | 1000    |


ROC-AUC Score: 0.7360719999999998



```
cm_dt = confusion_matrix(y_test, y_pred_dt)
disp_dt = ConfusionMatrixDisplay(confusion_matrix=cm_dt, display_labels=["Class 0", "Class 1"])
```

```
# Plot the confusion matrix
plt.figure(figsize=(8, 6))
disp_dt.plot(cmap="Blues", values_format="d")
plt.title("Confusion Matrix for Decision Tree")
plt.show()
```

&lt;Figure size 800x600 with 0 Axes&gt;



```
# Define the pipeline with Logistic Regression
pipeline_lr = Pipeline([
    ('logreg', LogisticRegression(solver='liblinear')) # 'liblinear' solver is suitable for binary classification
])

# Define the parameter grid for GridSearchCV
param_grid_lr = {
    'logreg__C': [0.01, 0.1, 1, 10, 100], # Regularization strength
    'logreg__penalty': ['l1', 'l2'] # L1 and L2 regularization
}

# Setup GridSearchCV
grid_lr = GridSearchCV(pipeline_lr, param_grid_lr, cv=5, scoring='accuracy', verbose=1)
# Train the model with GridSearchCV on the training data
grid_lr.fit(X_train, y_train) # assuming y_train contains the labels for training

# Best model, predict and evaluate
best_lr = grid_lr.best_estimator_
y_pred_lr = best_lr.predict(X_test)
y_prob_lr = best_lr.predict_proba(X_test)[:, 1] # Probability scores for ROC curve

# Display best parameters and cross-validation score
print("Best parameters:", grid_lr.best_params_)
print("Best cross-validation score: {:.2f}".format(grid_lr.best_score_))

# Print classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred_lr)) # assuming y_test contains the labels for testing

# Print ROC-AUC score
print("ROC-AUC Score:", roc_auc_score(y_test, y_prob_lr))


# Plot ROC curve
fpr_lr, tpr_lr, _ = roc_curve(y_test, y_prob_lr, pos_label=1)
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label='Logistic Regression (AUC = {:.2f})'.format(roc_auc_score(y_test, y_prob_lr)))
plt.plot([0, 1], [0, 1], 'r--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Logistic Regression')
plt.legend(loc="lower right")
plt.show()
```

↗ Fitting 5 folds for each of 10 candidates, totalling 50 fits  
 Best parameters: {'logreg\_\_C': 1, 'logreg\_\_penalty': 'l2'}  
 Best cross-validation score: 0.88

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.86      | 0.90   | 0.88     | 500     |
| 1            | 0.89      | 0.85   | 0.87     | 500     |
| accuracy     |           |        | 0.88     | 1000    |
| macro avg    | 0.88      | 0.88   | 0.88     | 1000    |
| weighted avg | 0.88      | 0.88   | 0.88     | 1000    |

ROC-AUC Score: 0.947784



```
cm_lr = confusion_matrix(y_test, y_pred_lr)
disp_lr = ConfusionMatrixDisplay(confusion_matrix=cm_lr, display_labels=["Class 0", "Class 1"])
```

```
# Plot the confusion matrix
plt.figure(figsize=(8, 6))
disp_lr.plot(cmap="Blues", values_format="d")
plt.title("Confusion Matrix for Logistic Regression")
plt.show()
```

↗ <Figure size 800x600 with 0 Axes>

Confusion Matrix for Logistic Regression